


Threads

(Operating Systems)




Deepika H V

C-DAC Bengaluru


(deepikahv@cdac.in)

Symmetric MultiProcessor

- **SMP – computer architecture where two or more identical processors can connect to a single shared memory.**



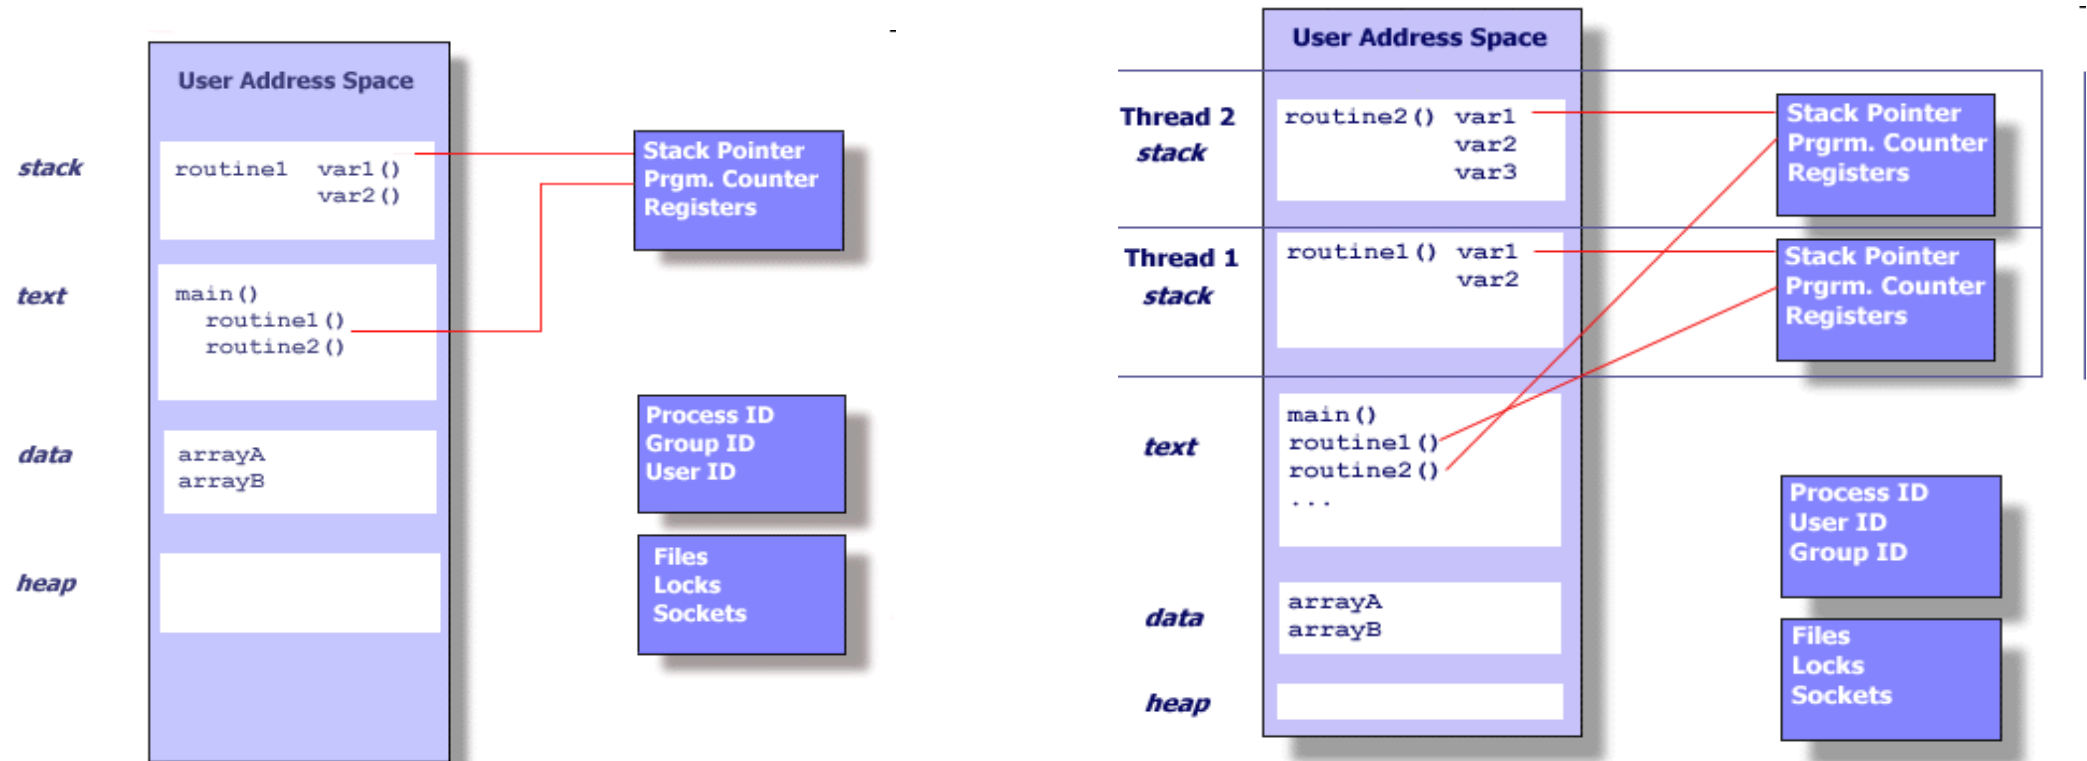
Shared Memory Machine



- **What is process?**
 - Program in execution
- **What is thread?**
 - Is an independent /different stream of control that can execute its instructions independently and can use the process resources
- **What is the connection b/w process and thread ?**

Thread

- Imagine a main program (a.out) that contains a number of procedures. Then imagine all of these procedures being able to run simultaneously and/or independently. That would describe a "multi-threaded" program



- **So, in summary, in the UNIX environment a thread:**
 - Exists within a process and uses the process resources
 - Has its own independent flow of control as long as its parent process exists and the OS supports it
 - Duplicates only the essential resources it needs to be independently schedulable
 - May share the process resources with other threads that act equally independently (and dependently)
 - Dies if the parent process dies
 - Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.

- **As threads within the same process share resources:**
 - Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads.
 - Two pointers having the same value point to the same data.
 - Reading and writing to the same memory locations is possible, and therefore requires explicit synchronization by the programmer.

	Process	Threads
1	Process is heavy weight or resource intensive.	Thread is light weight, taking lesser resources than a process.
2	Process switching needs interaction with operating system.	Thread switching does not need to interact with operating system.
3	If one process is blocked, then no other process can execute until the first process is unblocked.	While one thread is blocked and waiting, a second thread in the same task can run.
4	Multiple processes without using threads use more resources.	Multiple threaded processes use fewer resources.
5	In multiple processes each process operates independently of the others.	One thread can read, write or change another thread's data.
<div>Sep 27- Oct 4, 2021</div> <div>DAC , Operating System – Threads</div> <div>7</div>		

Value of Using Threads

- Performance gains from multiprocessing hardware (parallelism)
- Increased application throughput
- Increased application responsiveness
- Replacing process-to-process communications
- Efficient use of system resources
- Simplified signal handling
- The ability to make use of the inherent concurrency of distributed objects


Types of Threads

	User Level Thread	Kernel Level Threads
1	User-level threads are faster to create and manage.	Kernel-level threads are slower to create and manage.
2	Implementation is by a thread library at the user level.	Operating system supports creation of Kernel threads.
3	User-level thread is generic and can run on any operating system.	Kernel-level thread is specific to the operating system.

Multithreading Models


- **One to One**

- Windows NT, Linux




- **Many to one**

- Solaris Green Threads and GNU portable threads



- **Many to Many**

- Solaris prior to ver 9
- Windows 2000



Posix Threads (pthreads)

- specified by the IEEE POSIX 1003.1c standard (1995).
- set of C programming types & procedure calls, implemented with a pthread.h header file and a thread library.

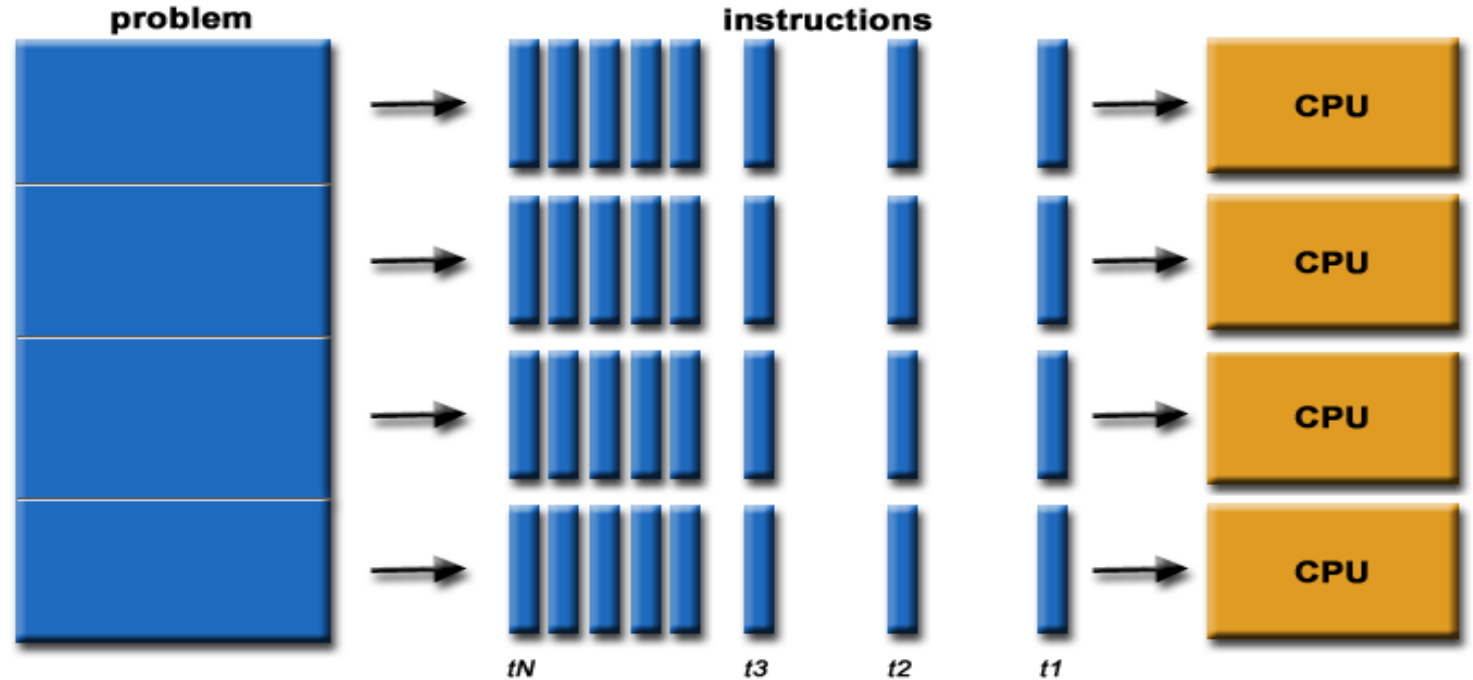
- **Why Pthreads.**

- 5000 threads/process.

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
AMD 2.4 GHz Opteron (8cpus/node)	41.07	60.08	9.01	0.66	0.19	0.43
IBM 1.9 GHz POWER5 p5-575 (8cpus/node)	64.24	30.78	27.68	1.75	0.69	1.10
IBM 1.5 GHz POWER4 (8cpus/node)	104.05	48.64	47.21	2.01	1.00	1.52
INTEL 2.4 GHz Xeon (2 cpus/node)	54.95	1.54	20.78	1.64	0.67	0.90
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.54	1.07	22.22	2.03	1.26	0.67

When is threading Useful

- Independent Tasks
- Servers
- Repetitive tasks
- Asynchronous events



- **Considerations For Thread Programming**

- Problem partitioning and complexity
- Load balancing
- Data dependencies
- Synchronization and race conditions
- Data communications
- Memory, I/O issues

Naming convention

Routine Prefix	Functional Group
pthread_	Threads themselves and miscellaneous subroutines
pthread_attr_	Thread attributes objects
pthread_mutex_	Mutexes
pthread_mutexattr_	Mutex attributes objects.
pthread_cond_	Condition variables
pthread_condattr_	Condition attributes objects
pthread_key_	Thread-specific data keys

Compilation

Compiler / Platform	Compiler Command	Description
IBM AIX	<code>xlc_r / cc_r</code>	C (ANSI / non-ANSI)
	<code>xlc_r</code>	C++
	<code>xlf_r -qnosave</code> <code>xlf90_r -qnosave</code>	Fortran - using IBM's Pthreads API (non-portable)
INTEL Linux	<code>icc -pthread</code>	C
	<code>icpc -pthread</code>	C++
PathScale Linux	<code>pathcc -pthread</code>	C
	<code>pathCC -pthread</code>	C++
PGI Linux	<code>pgcc -lpthread</code>	C
	<code>pgCC -lpthread</code>	C++
GNU Linux, AIX	<code>gcc -pthread</code>	GNU C
	<code>g++ -pthread</code>	GNU C++

Concept

- Concept of opaque objects pervades the design of API.
- Pthreads has over 100 subroutines
- For portability, pthread.h header file should be used for accessing pthread library.
- POSIX standard defined only for C language
- Once threads are created they are peers and may create other threads.
- Maximum number of threads created is implementation dependent.

1. Thread Management

- `pthread_create (thread,attr,start_routine,arg)`
- `pthread_exit (status)`
- `pthread_attr_init (attr)`
- `pthread_attr_destroy (attr)`
- `pthread_join (threadid,status)`
- `pthread_detach (threadid,status)`

Pthread_create

- **pthread_create (thread, attr, start_routine, arg)**
 - creates a new thread and makes it executable.
- **thread**: An unique identifier for the new thread returned by the subroutine.
- **attr**: An attribute object that may be used to set thread attributes. NULL for the default values.
- **start_routine**: the C routine that the thread will execute once it is created.
- **arg**: A single argument that may be passed to start_routine. It must be **passed by reference as a pointer cast of type void**. NULL may be used if no argument is to be passed.

Termination

- Thread returns from main routine.
- Thread calls **pthread_exit (status)** . This is used to explicitly exit a thread
- the pthread_exit() routine **does not close files**; any files opened inside the thread will remain open after the thread is terminated.
- Thread is cancelled by other thread – **pthread_cancel()**
- Entire process is terminated.

Example Code - Pthread Creation and Termination

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS      5

void *PrintHello(void *threadid)
{
    long tid;
    tid = (long)threadid;
    printf("Hello World! It's me, thread %ld!\n", tid);
    pthread_exit(NULL);
}

int main (int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc;
    long t;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```



Example 2 - Thread Argument Passing

This example shows how to setup/pass multiple arguments via a structure. Each thread receives a unique instance of the structure.

```
struct thread_data{
    int  thread_id;
    int  sum;
    char *message;
};

struct thread_data thread_data_array[NUM_THREADS];

void *PrintHello(void *threadarg)
{
    struct thread_data *my_data;
    ...
    my_data = (struct thread_data *) threadarg;
    taskid = my_data->thread_id;
    sum = my_data->sum;
    hello_msg = my_data->message;
    ...
}

int main (int argc, char *argv[])
{
    ...
    thread_data_array[t].thread_id = t;
    thread_data_array[t].sum = sum;
    thread_data_array[t].message = messages[t];
    rc = pthread_create(&threads[t], NULL, PrintHello,
        (void *) &thread_data_array[t]);
    ...
}
```

- `pthread_attr_getstacksize (attr, stacksize)`
- `pthread_attr_setstacksize (attr, stacksize)`
- `pthread_attr_getstackaddr (attr, stackaddr)`
- `pthread_attr_setstackaddr (attr, stackaddr)`

Mutex

- **Mutex** is an abbreviation for "**mutual exclusion**". **Mutex** variables are one of the primary means of implementing thread synchronization and for protecting shared data when multiple writes occur.
- **Mutexes** can be used to **prevent "race" conditions**.

Thread 1	Thread 2	Balance
Read balance: \$1000		\$1000
	Read balance: \$1000	\$1000
	Deposit \$200	\$1000
Deposit \$200		\$1000
Update balance \$1000+\$200		\$1200
	Update balance \$1000+\$200	\$1200

Sequence

- Create and initialize a mutex variable
- Several threads attempt to lock the mutex
- only one succeeds and that thread owns the mutex
- The owner thread performs some set of actions
- The owner unlocks the mutex
- Another thread acquires the mutex and repeats the process
- Finally the mutex is destroyed

Mutex Routines

- `pthread_mutex_init (mutex,attr)`
- `pthread_mutex_destroy (mutex)`
- `pthread_mutexattr_init (attr)`
- `pthread_mutexattr_destroy (attr)`
- `pthread_mutex_lock (mutex)`
- `pthread_mutex_trylock (mutex)`
- `pthread_mutex_unlock (mutex)`

- Condition variables provide yet another way for threads to synchronize. While mutexes implement **synchronization by controlling thread access to data**, condition variables allow threads to synchronize based upon the **actual value of data**.
- Without condition variables, the programmer would need to have **threads continually polling** (possibly in a critical section), to check if the condition is met. This can be very resource consuming since the thread would be continuously busy in this activity. A condition variable is a way to achieve the same goal without polling.
- A condition variable is always used in **conjunction with a mutex lock**.

Main Thread

- Declare and initialize global data/variables which require synchronization (such as "count")
- Declare and initialize a condition variable object
- Declare and initialize an associated mutex
- Create threads A and B to do work

Thread A

- Do work up to the point where a certain condition must occur (such as "count" must reach a specified value)
- Lock associated mutex and check value of a global variable
- Call `pthread_cond_wait()` to perform a blocking wait for signal from Thread-B. Note that a call to `pthread_cond_wait()` automatically and atomically unlocks the associated mutex variable so that it can be used by Thread-B.
- When signalled, wake up. Mutex is automatically and atomically locked.
- Explicitly unlock mutex
- Continue

Thread B

- Do work
- Lock associated mutex
- Change the value of the global variable that Thread-A is waiting upon.
- Check value of the global Thread-A wait variable. If it fulfills the desired condition, signal Thread-A.
- Unlock mutex.
- Continue

Main Thread

Join / Continue

- `pthread_cond_init (condition,attr)`
- `pthread_cond_destroy (condition)`
- `pthread_condattr_init (attr)`
- `pthread_condattr_destroy (attr)`
- `pthread_cond_wait (condition,mutex)`
- `pthread_cond_signal (condition)`
- `pthread_cond_broadcast (condition)`

Thank You